

Cedar Pilot Field Guide

This is a fictional, self-contained document created for testing RAGGlass. All Cedar project facts are invented fixture data, not measurements of a real deployment.

1. Pilot overview

The Cedar pilot is maintained by Mira Chen. The pilot begins on September 1, 2026 and ends on September 30, 2026. Its library contains 12 native-text PDF documents.

The first release accepts PDFs with selectable native text. Optical character recognition (OCR) is disabled. Scanned image-only PDFs are outside this release's scope.

The pilot is intended for engineers inspecting the path from an original PDF to parsed content, retrieved evidence, and a grounded answer. The application runs locally.

Data ownership

Original PDFs and parsed exports are retained on local disk. SQLite stores document metadata and query history. Qdrant stores the persistent vector index.

This PDF and its generation script are dedicated to the public domain under CC0 1.0. You may copy, modify, and redistribute this fixture.

Limits and defaults

The following table defines the Cedar pilot's configuration. These are fixture specifications, not benchmark results. Token counts refer to the embedding tokenizer.

2. Configuration table

Setting	Value	Meaning
Maximum upload size	30 MB	Per PDF document
Maximum document length	200 pages	Per PDF document
Chunk window	240 tokens	Page-aware token window
Chunk overlap	32 tokens	Between adjacent windows
Retrieved chunks (Top K)	5	Maximum chunks per query
Minimum cosine score	0.70	Retrieval score threshold

For example, one PDF larger than 30 MB must be split or reduced before upload. The page limit is separate from the file size limit.

The default retrieval returns at most five chunks; fewer may pass the minimum score. A cosine similarity score is not a calibrated probability of correctness.

Trace every answer

3. Citation contract

A model cites retrieved chunk IDs. The backend validates that every cited chunk belongs to the current query's evidence, then maps it to the original document and page numbers.

If the evidence does not answer the question, the answer must say that it cannot be confirmed from the document. Unsupported sources must not be shown as valid citations.

A chunk retains its document ID, document SHA-256 hash, chunk ID, page numbers, parsing settings, and source coordinates when available. Missing coordinates are explicitly marked.

Reproducing a query

Every query saves the model name, model HTTP endpoint, prompt, embedding model revision, chunk settings, retrieval settings, evidence, validated citations, and measured stage timings.

After restarting the application, engineers can reopen existing documents and query history. The original files, SQLite database, and Qdrant storage must be kept together.

Information deliberately omitted

This guide does not specify the pilot's annual electricity cost, GPU purchase price, passwords, or any benchmark latency. Questions about these values cannot be answered from this document.